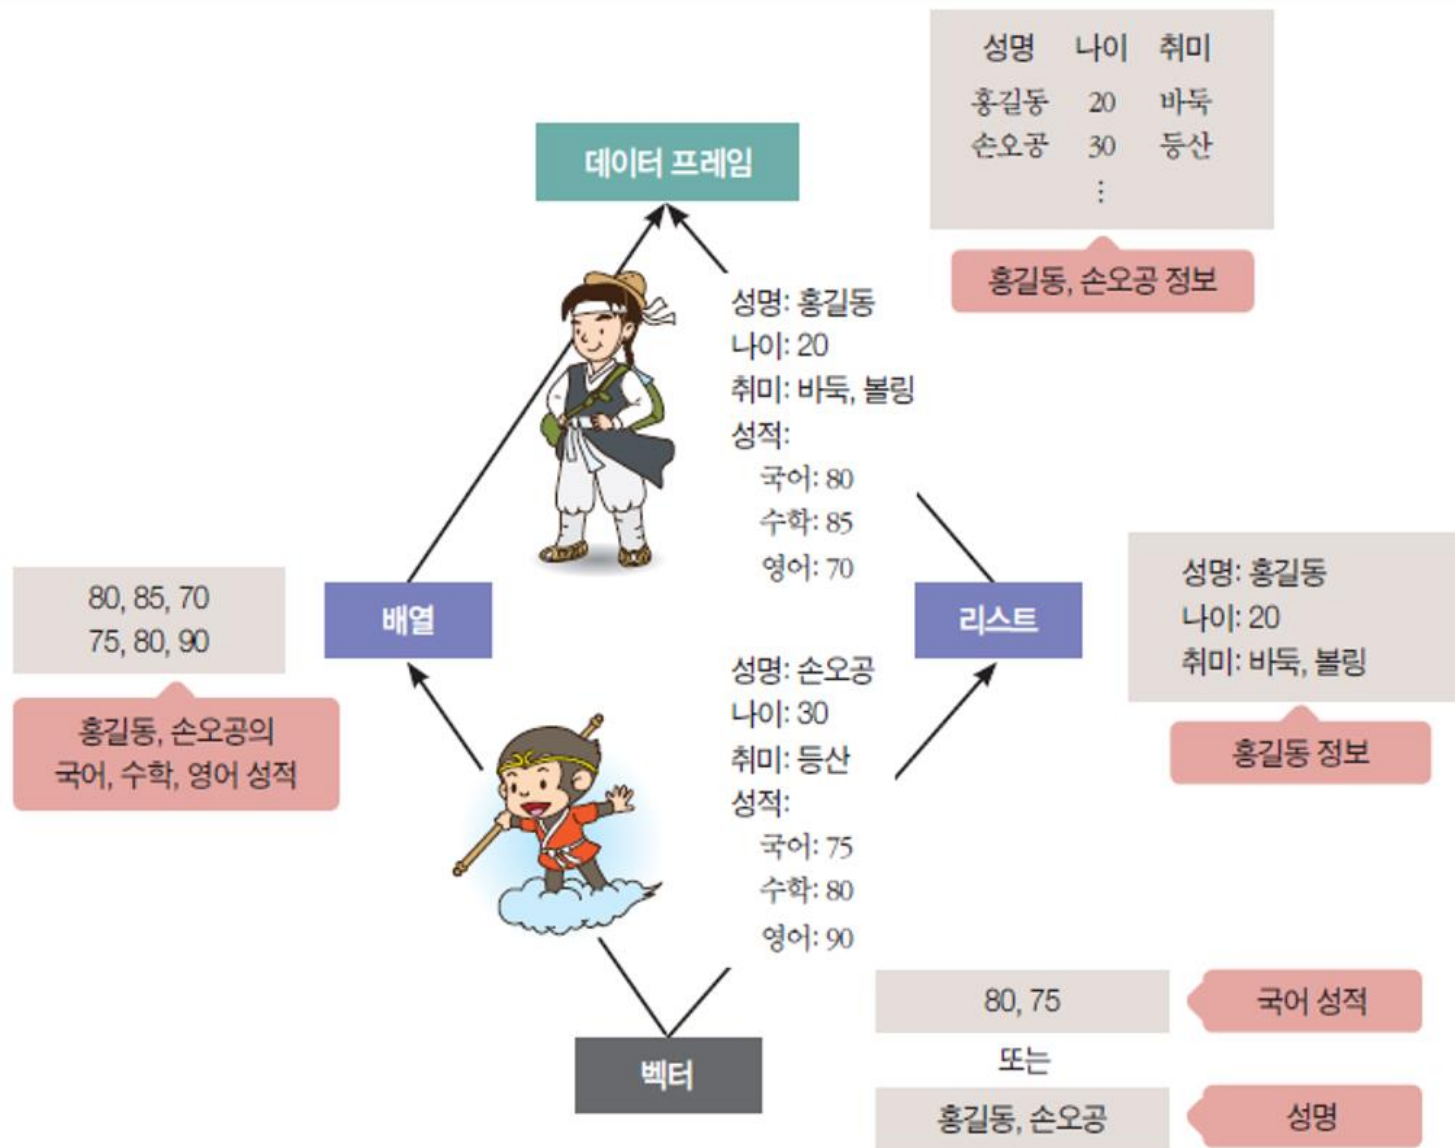


# Vectors

Bok, Jong Soon  
javaexpert@nate.com  
[www.javaexpert.info](http://www.javaexpert.info)

# Data Types in R



# Vectors

- The fundamental data type in R.
- Scalar : single-number variables.
- In C-language family.  

```
int x;  
int y[3];
```
- But in R, numbers are actually considered one-element vectors.
- R variable types are called *modes*.
- All elements in a vector must have the same mode.
- It can be integer, numeric (floating-point number), character (string), logical (Boolean), complex, and so on.

```
1  
2 kor <- 85  
3 name <- "홍길동"  
4 grade <- "B"  
5  
6 kor  
7 name  
8 grade  
9 |
```



9:1 (Top Level) ↕

Console C:/R Home/ ↗

```
> name <- "홍길동"  
> grade <- "B"  
> kor  
[1] 85  
> name  
[1] "홍길동"  
> grade  
[1] "B"  
> |
```


# Single Element Vector Creation

- Even when you write just one value in R, it becomes a vector of length 1 and belongs to one of the above vector types.

```
Console ~/  
> # Atomic vector of type character.
> print("abc")
[1] "abc"
>
> # Atomic vector of type double.
> print(12.5)
[1] 12.5
>
> # Atomic vector of type integer.
> print(63L)
[1] 63
>
> # Atomic vector of type logical.
> print(TRUE)
[1] TRUE
>
> # Atomic vector of type complex.
> print(2 + 3i)
[1] 2+3i
>
> # Atomic vector of type raw.
> print(charToRaw('hello'))
[1] 68 65 6c 6c 6f
```

# Multiple Elements Vector Creation

- Using colon operator( : ) with numeric data.

```
Console ~/   
> # Creating a sequence from 5 to 13.  
> v <- 5:13  
> print(v)  
[1] 5 6 7 8 9 10 11 12 13  
>  
> # Creating a sequence from 6.6 to 12.6  
> v <- 6.6 : 12.6  
> print(v)  
[1] 6.6 7.6 8.6 9.6 10.6 11.6 12.6  
>  
> # If the final element specified does not belong to the sequence then it is discarded.  
> v <- 3.8 : 11.4  
> print(v)  
[1] 3.8 4.8 5.8 6.8 7.8 8.8 9.8 10.8
```

# Multiple Elements Vector Creation (Cont.)

- Using `seq()` function.

Console ~/ ↗

```
> # Create vector with elements from 5 to 9 incrementing by 0.4.  
> print(seq(5, 9, by = 0.4))  
[1] 5.0 5.4 5.8 6.2 6.6 7.0 7.4 7.8 8.2 8.6 9.0  
>
```

# Multiple Elements Vector Creation (Cont.)

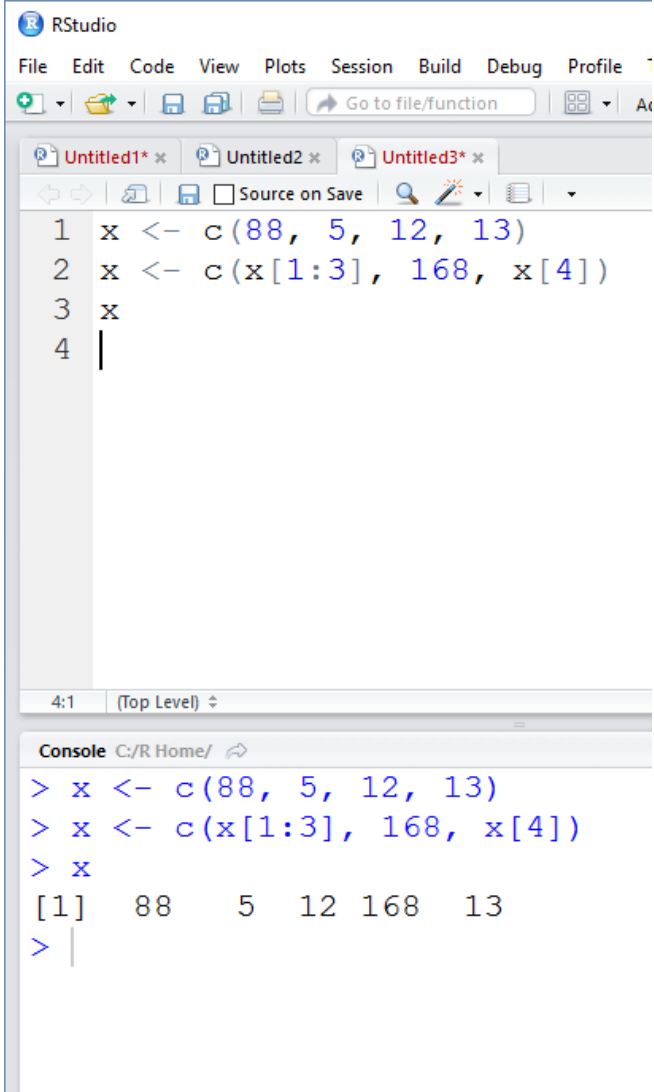
- Using the `c()` function.
- The non-character values are coerced to character type if one of the elements is a character.

Console ~/ ↗

```
> # The logical and numeric values are converted to characters.  
> s <- c('apple', 'red', 5, TRUE)  
> print(s)  
[1] "apple" "red"   "5"     "TRUE"  
>
```

# Adding and Deleting Vector Elements

- Vectors are stored like arrays in C.
- But, cannot insert or delete elements.
- The size of a vector is determined at its creation.
- Unlike vector indices in ALGOL-family languages, such as C and Python, vector indices in R begin at **1**.



The screenshot shows the RStudio interface. The editor window contains the following R code:

```
1 x <- c(88, 5, 12, 13)
2 x <- c(x[1:3], 168, x[4])
3 x
4 |
```

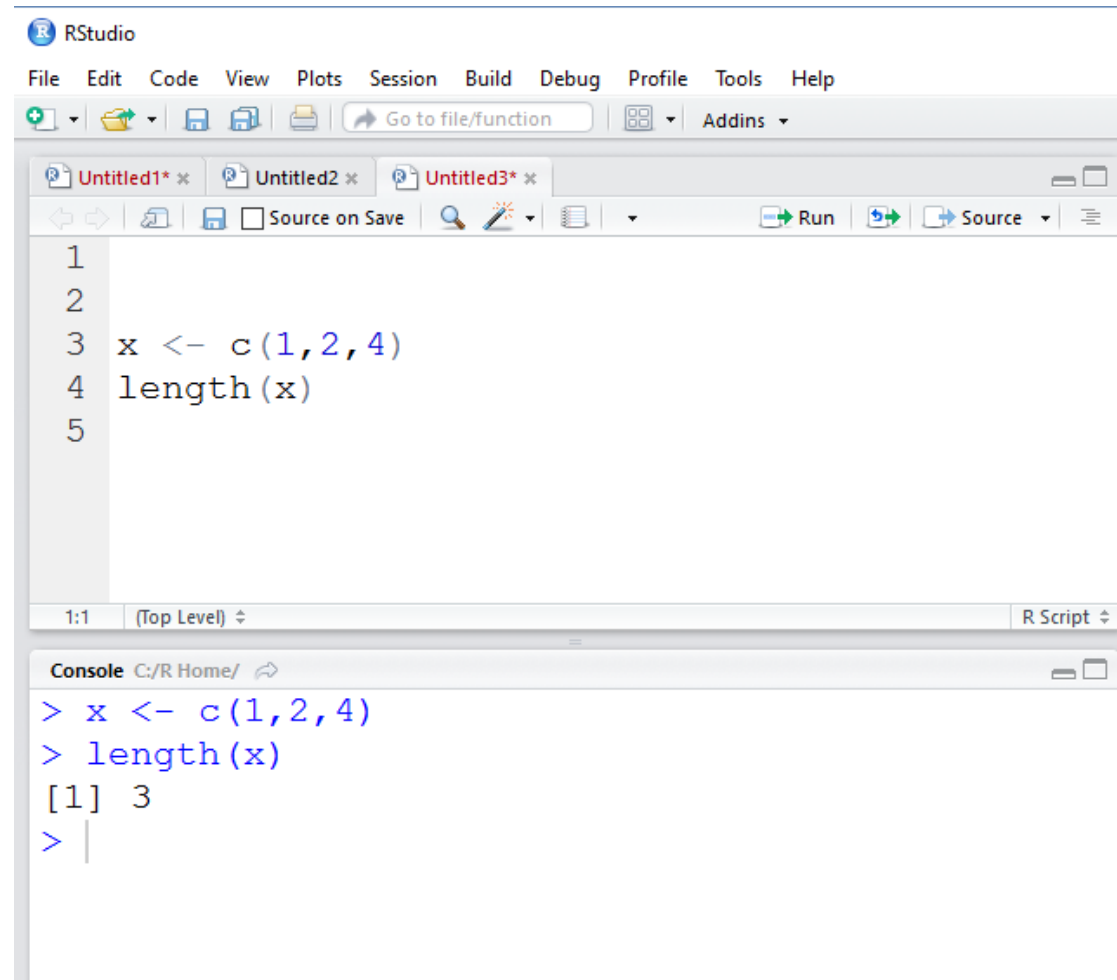
The console window shows the execution of the code:

```
> x <- c(88, 5, 12, 13)
> x <- c(x[1:3], 168, x[4])
> x
[1] 88 5 12 168 13
> |
```



# Obtaining the Length of a Vector

- You can obtain the length of a vector by using the **length()** function.



The screenshot shows the RStudio interface. The script editor contains the following code:

```
1  
2  
3 x <- c(1, 2, 4)  
4 length(x)  
5
```

The console shows the execution of the code:

```
> x <- c(1, 2, 4)  
> length(x)  
[1] 3  
> |
```

# Matrices and Arrays as Vectors

- Arrays and are actually vectors too.
- Matrices have the number of rows and columns.

```
1 m <- matrix(c(1,2,3,4), nrow=2)
2 m
3 m + 10:13
4 |
```

4:1 (Top Level) ↕

Console C:/R Home/ ↗

```
> m <- matrix(c(1,2,3,4), nrow=2)
> m
      [,1] [,2]
[1,]    1    3
[2,]    2    4
> m + 10:13
      [,1] [,2]
[1,]   11   15
[2,]   13   17
> |
```

# Declarations

- Unlike C, Instead array, must create **y** first, for instance this way:

```
> y <- vector(length=2)
```

```
> y[1] <- 5
```

```
> y[2] <- 12
```

- The following will also work:

```
> y <- c(5,12)
```

## Declarations (Cont.)

- The following sequence of events is perfectly valid.

```
> x <- c(1,5)
```

```
> x
```


```
[1] 1 5
```

```
> x <- "abc"
```

- First, **x** is associated with a numeric vector, then with a string.

# Recycling

- When applying an operation to two vectors.
- Requires them to be the same length, R automatically recycles, or repeats, the shorter one, until it is long enough to match the longer one.
- Here is an example:

```
Console C:/R Home/   
> c(1,2,4) + c(6,0,9,20,22)  
[1] 7 2 13 21 24  
Warning message:  
In c(1, 2, 4) + c(6, 0, 9, 20, 22) :  
  longer object length is not a multiple of shorter  
  object length  
> |
```

## Recycling (Cont.)

- The shorter vector was recycled, so the operation was taken to be as follows:

```
> c(1,2,4,1,2) + c(6,0,9,20,22)
```

# Vector Arithmetic and Logical Operations

- R is a functional language.
- Every operator, including **+** is actually a function.

```
> 2+3
```

```
[1] 5
```

```
> "+"(2,3)
```

```
[1] 5
```

# Vector Arithmetic and Logical Operations (Cont.)

- Can add vectors, and the **+** operation will be applied element-wise.

```
> x <- c(1, 2, 4)
> x + c(5, 0, -1)
[1] 6 2 3
```

- If you are familiar with linear algebra, you may be surprised at what happens when we multiply two vectors.

```
> x * c(5, 0, -1)
[1] 5 0 -4
```



# Vector Arithmetic and Logical Operations (Cont.)

- The same principle applies to other numeric operators.
- Here's an example:

```
> x <- c(1,2,4)
```

```
> x / c(5,4,-1)
```

```
[1] 0.2 0.5 -4.0
```

```
> x %% c(5,4,-1)
```

```
[1] 1 2 0
```

# Vector Indexing

- One of the most important and frequently used operations in R is that of indexing vectors.
- Form a subvector by picking elements of the given vector for specific indices.

```
> y <- c(1.2, 3.9, 0.4, 0.12)
> y[c(1,3)] # extract elements 1 and 3 of y
[1] 1.2 0.4
> y[2:3]
[1] 3.9 0.4
> v <- 3:4
> y[v]
[1] 0.40 0.12
```

# Vector Indexing (Cont.)

- Note that duplicates are allowed.

```
> x <- c(4, 2, 17, 5)
> y <- x[c(1, 1, 3)]
> y
[1] 4 4 17
```

- Negative subscripts mean that we want to exclude the given elements in our output.

```
> z <- c(5, 12, 13)
> z[-1]      # exclude element 1
[1] 12 13
> z[-1:-2]   # exclude elements 1 through 2
[1] 13
```

## Vector Indexing (Cont.)

- It is often useful to use the `length()` function.
- The following code will do just that:

```
> z <- c(5,12,13)
> z[1:(length(z)-1)]
[1] 5 12
```

- Or more simply:

```
> z[-length(z)]
[1] 5 12
```

- This is more general than using `z[1:2]`.

# Generating Useful Vectors with the : Operator

- There are a few R operators that are especially useful for creating vectors.
- Let's start with the colon operator `:`.
- It produces a vector consisting of a range of numbers.

```
> 5:8
```

```
[1] 5 6 7 8
```

```
> 5:1
```

```
[1] 5 4 3 2 1
```

## Generating Useful Vectors with the : Operator (Cont.)

- You may recall that it was used earlier in this chapter in a loop context, as follows:

```
for (i in 1:length(x)) {
```

- Beware of operator precedence issues.

```
> i <- 2
```

```
> 1:i-1 # this means (1:i) - 1, not 1:(i-1)
```

```
[1] 0 1
```

```
> 1:(i-1)
```

```
[1] 1
```

# Generating Vector Sequences with seq()

- A generalization of : is the `seq()` (or sequence) function.
- Generates a sequence in arithmetic progression.

```
> seq(from=12, to=30, by=3)
[1] 12 15 18 21 24 27 30
```
- The spacing can be a non-integer value, too, say 0.1.

```
> seq(from=1.1, to=2, length=10)
[1] 1.1 1.2 1.3 1.4 1.5 1.6 1.7 1.8 1.9
2.0
```

# Generating Vector Sequences with seq() (Cont.)

Console C:/R Home/ ↗

```
>
> x <- c(5, 12, 13)
> x
[1] 5 12 13
> seq(x)
[1] 1 2 3
> x <- NULL
> x
NULL
> seq(x)
integer(0)
> |
```

- You can see that **seq(x)** gives us the same result as **1:length(x)** if **x** is not empty, but it correctly evaluates to **NULL** if **x** is empty, resulting in zero iterations in the above loop.



# Repeating Vector Constants with rep()

- The `rep()` (or `repeat`) function allows us to conveniently put the same constant into long vectors.
- The call form is `rep(x, times)`.
- Creates a vector of `times * length(x)` elements
- That is, `times` copies of `x`.

```
> x <- rep(8, 4)
```

```
> x
```

```
[1] 8 8 8 8
```

```
> rep(c(5, 12, 13), 3)
```

```
[1] 5 12 13 5 12 13 5 12 13
```

```
> rep(1:3, 2)
```

```
[1] 1 2 3 1 2 3
```

## Repeating Vector Constants with rep() (Cont.)

- There is a named argument `each`, with very different behavior, which interleaves the copies of **x**.

```
> rep(c(5,12,13),each=2)
[1] 5 5 12 12 13 13
```

# Using all() and any()

- The **any()** and **all()** functions are handy shortcuts.
- They report whether any or all of their arguments are **TRUE**.

```
> x <- 1:10
```

```
> any(x > 8)
```

```
[1] TRUE
```

```
> any(x > 88)
```

```
[1] FALSE
```

```
> all(x > 88)
```

```
[1] FALSE
```

```
> all(x > 0)
```

```
[1] TRUE
```

## Using all() and any() (Cont.)

- For example, suppose that R executes the following:  
`> any(x > 8)`
- It first evaluates `x > 8`, yielding this:  
  
`(FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, TRUE, TRUE)`
- The `any()` function then reports whether any of those values is `TRUE`.
- The `all()` function works similarly and reports if all of the values are `TRUE`.

# Vector In, Vector Out

- You saw examples of vectorized functions earlier in the chapter, with the **+** and **\*** operators.
- Another example is **>**.

```
> u <- c(5, 2, 8)
```

```
> v <- c(1, 3, 9)
```

```
> u > v
```

```
[1] TRUE FALSE FALSE
```

## Vector In, Vector Out (Cont.)

- An R function uses *vectorized* operations.
- Here is an example:

```
> w <- function(x) return(x+1)
```

```
> w(u)
```

```
[1] 6 3 9
```

- Here, `w()` uses `+`, which is *vectorized*, so `w()` is *vectorized* as well.

## Vector In, Vector Out (Cont.)

```
> sqrt(1:9)
[1] 1.000000 1.414214 1.732051 2.000000 2.236068
2.449490 2.645751 2.828427
[9] 3.000000
```

```
> y <- c(1.2, 3.9, 0.4)
> z <- round(y)
> z
[1] 1 4 0
```

- The point is that the `round()` function is applied individually to each element in the vector `y`.

## Vector In, Vector Out (Cont.)

- As mentioned earlier, even operators such as **+** are really functions.
- For example, consider this code:

```
> y <- c(12, 5, 13)
```

```
> y + 4
```

```
[1] 16 9 17
```

- The reason element-wise addition of 4 works here is that the **+** is actually a function!

```
> '+'(y, 4)
```

```
[1] 16 9 17
```



# Vector In, Vector Out (Cont.)

Console C:/R Home/ ↗

```
>
> f <- function(x, c) {
+   return ((x + c) ^ 2)
+ }
> f
function(x, c) {
  return ((x + c) ^ 2)
}
> f(1:3, 0)
[1] 1 4 9
> f(1:3, 1)
[1] 4 9 16
>
>
> f(1:3, 1:3)
[1] 4 16 36
>
```

# Vector In, Matrix Out

- What if our function itself is vector-valued, as `z12()` is here:

```
z12 <- function(z) return(c(z, z^2))
```

- Applying `z12()` to 5, say, gives us the two-element vector `(5, 25)`.
- If we apply this function to an eight-element vector, it produces 16 numbers:

```
x <- 1:8
```

```
> z12(x)
```

```
[1] 1 2 3 4 5 6 7 8 1 4 9 16 25 36 49 64
```

## Vector In, Matrix Out (Cont.)

- It might be more natural to have these arranged as an 8-by-2 matrix, which we can do with the matrix function:

```
> matrix(z12(x), ncol=2)
```

```
[,1] [,2]
```

```
[1,] 1  1
```

```
[2,] 2  4
```

```
[3,] 3  9
```

```
[4,] 4 16
```

```
[5,] 5 25
```

```
[6,] 6 36
```

```
[7,] 7 49
```

```
[8,] 8 64
```

# Using NA

- In statistical data sets, often encounter missing data, which we represent in R with the value **NA**.

```
> x <- c(88, NA, 12, 168, 13)
```

```
> x
```

```
[1] 88 NA 12 168 13
```

```
> mean(x)
```

```
[1] NA
```

```
> mean(x, na.rm=T)
```

```
[1] 70.25
```

```
> x <- c(88, NULL, 12, 168, 13)
```

```
> # R automatically skipped over the NULL value
```

```
> mean(x)
```

```
[1] 70.25
```

## Using NA (Cont.)

- There are multiple **NA** values, one for each mode:

```
Console C:/R Home/ ↗
>
> x <- c(5, NA, 12)
> mode(x[1])
[1] "numeric"
> mode(x[2])
[1] "numeric"
> y <- c("abc", "def", NA)
> mode(y[2])
[1] "character"
> mode(y[3])
[1] "character"
>
```

# Using NULL

- One use of **NULL** is to build up vectors in loops.
- Each iteration adds another element to the vector.
- In this simple example, we build up a vector of even numbers:

```
# build up a vector of the even numbers in 1:10
```

```
> z <- NULL
```

```
> for (i in 1:10) if (i %% 2 == 0) z <- c(z,i)
```

```
> z
```

```
[1] 2 4 6 8 10
```

## Using NULL (Cont.)

- If we were to use **NA** instead of **NULL** in the preceding example, we would pick up an unwanted **NA**:

```
> z <- NA
> for (i in 1:10) if (i %% 2 == 0) z <-
c(z,i)
> z
[1] NA 2 4 6 8 10
```

## Using NULL (Cont.)

- **NULL** values really are counted as nonexistent, as you can see here:

```
> u <- NULL
```

```
> length(u)
```

```
[1] 0
```

```
> v <- NA
```

```
> length(v)
```

```
[1] 1
```

- **NULL** is a special R object with no mode.



# Generating Filtering Indices

- Let's start with a simple example:

```
> z <- c(5, 2, -3, 8)
```

```
> w <- z[z * z > 8]
```

```
> w
```

```
[1] 5 -3 8
```

## Generating Filtering Indices (Cont.)

- Let's look at it done piece by piece:

```
> z <- c(5, 2, -3, 8)
```

```
> z
```

```
[1] 5 2 -3 8
```

```
> z * z > 8
```

```
[1] TRUE FALSE TRUE TRUE
```

## Generating Filtering Indices (Cont.)

- The operator **>**, like **+**, is actually a function.
- Let's look at an example of that last point:

```
> ">"(2, 1)
```

```
[1] TRUE
```

```
> ">"(2, 5)
```

```
[1] FALSE
```

# Generating Filtering Indices (Cont.)

- Thus, the following:

`z*z > 8`

- is really this:

`">"(z*z, 8)`

- The resulting Boolean values are used to cull out the desired elements of z:

```
> z <- c(5, 2, -3, 8)
```

```
> z
```

```
[1]  5  2 -3  8
```

```
> z[c(TRUE, FALSE, TRUE, TRUE)]
```

```
[1]  5 -3  8
```

# Generating Filtering Indices (Cont.)

- Let's look at this code.

```
> z <- c(5, 2, -3, 8)
> j <- z*z > 8
> j
[1] TRUE FALSE TRUE
TRUE
> y <- c(1, 2, 30, 5)
> y[j]
[1] 1 30 5
```

- Or, more compactly, we could write the following:

```
> z <- c(5, 2, -3, 8)
> y <- c(1, 2, 30, 5)
> y[z*z > 8]
[1] 1 30 5
```

# Filtering with the subset() Function

- When applied to vectors, the difference between using this function and ordinary filtering lies in the manner in which **NA** values are handled.

```
> x <- c(6, 1:3, NA, 12)
```

```
> x
```

```
[1] 6 1 2 3 NA 12
```

```
> x[x > 5]
```

```
[1] 6 NA 12
```

```
> subset(x, x > 5)
```

```
[1] 6 12
```

# The Selection Function `which()`

- We may just want to find the positions within `z` at which the condition occurs.
- We can do this using `which()`, as follows:

```
> z <- c(5, 2, -3, 8)
```

```
> which(z * z > 8)
```

```
[1] 1 3 4
```

# A Vectorized if-then-else: The ifelse() Function

- The usual if-then-else construct found in most languages, R also includes a *vectorized* version, the `ifelse()` function.
- The form is as follows:  
`ifelse(b,u,v)`
- Where `b` is a Boolean vector, and `u` and `v` are vectors.
- The return value is itself a vector; element `i` is `u[i]` if `b[i]` is `true`, or `v[i]` if `b[i]` is `false`.



## A Vectorized if-then-else: The ifelse() Function (Cont.)

- The concept is pretty abstract, so let's go right to an example:

```
> x <- 1:10
```

```
> y <- ifelse(x %% 2 == 0, 5, 12)
```

```
# %% is the mod operator
```

```
> y
```

```
[1] 12 5 12 5 12 5 12 5 12 5
```

## A Vectorized if-then-else: The ifelse() Function (Cont.)

- Here is another example:

```
> x <- c(5, 2, 9, 12)
```

```
> ifelse(x > 6, 2 * x, 3 * x)
```

```
[1] 15 6 18 24
```

# Testing Vector Equality

- Suppose we wish to test whether two vectors are equal.
- The naive approach, using `==`, won't work.

```
> x <- 1:3
```

```
> y <- c(1,3,4)
```

```
> x == y
```

```
[1] TRUE FALSE FALSE
```

## Testing Vector Equality (Cont.)

- Just like almost anything else in R, `==` is a function.

```
> "==" (3, 2)
```

```
[1] FALSE
```

```
> i <- 2
```

```
> "==" (i, 2)
```

```
[1] TRUE
```

## Testing Vector Equality (Cont.)

- One option is to work with the *vectorized* nature of `==`, applying the function `all()`:

```
> x <- 1:3
```

```
> y <- c(1,3,4)
```

```
> x == y
```

```
[1] TRUE FALSE FALSE
```

```
> all(x == y)
```

```
[1] FALSE
```

## Testing Vector Equality (Cont.)

- Or even better, we can simply use the identical function, like this:

```
> x <- 1:3  
> y <- c(1,3,4)  
> identical(x, y)  
[1] FALSE
```

# Testing Vector Equality (Cont.)

- Consider this little R session:

```
> x <- 1:2
> y <- c(1,2)
> x
[1] 1 2
> y
[1] 1 2
> identical(x,y)
[1] FALSE
> typeof(x)
[1] "integer"
> typeof(y)
[1] "double"
```

So, `:` produces integers while `c()` produces floating-point numbers.

# Vector Element Names

- We can assign or query vector element names via the **names()** function:

```
> x <- c(1, 2, 4)
```

```
> names(x)
```

```
NULL
```

```
> names(x) <- c("a", "b", "ab")
```

```
> names(x)
```

```
[1] "a" "b" "ab"
```

```
> x
```

```
a b ab
```

```
1 2 4
```



## Vector Element Names (Cont.)

- We can remove the names from a vector by assigning **NULL**:

```
> names(x) <- NULL  
> x  
[1] 1 2 4
```

- We can even reference elements of the vector by name:

```
> x <- c(1, 2, 4)  
> names(x) <- c("a", "b", "ab")  
> x["b"]  
b 2
```