

Matrices and Arrays

Bok, Jong Soon
javaexpert@nate.com
www.javaexpert.info

Matrices

- Are a vector with two additional attributes : the number of rows and the number of columns.
- Since matrices are vectors, they also have modes, such as numeric and character.

Arrays

- Can be multidimensional.
- For example, a three-dimensional array would consist of rows, columns, and layers, not just rows and columns as in the matrix case.

Creating Matrices


- **matrix(data, nrow, ncol, byrow, dimnames)**
 - **data** is the input vector which becomes the data elements of the matrix.
 - **nrow** is the number of rows to be created.
 - **ncol** is the number of columns to be created.
 - **byrow** is a logical clue. If **TRUE** then the input vector elements are arranged by row.
 - **dimname** is the names assigned to the rows and columns.

Creating Matrices (Cont.)

- Matrix row and column subscripts begin with **1**.
- For example, the upper-left corner of the matrix **a** is denoted **a[1,1]**.
- The internal storage of a matrix is in *column-major order*, meaning that first all of column 1 is stored, then all of column 2, and so on.

Creating Matrices (Cont.)

- One way to create a matrix is by using the **matrix()** function:

```
Console ~/ 
> y <- matrix(c(1, 2, 3, 4), nrow=2, ncol=2)
> y
      [,1] [,2]
[1,]    1    3
[2,]    2    4
>
>
>
>
> y <- matrix(c(1, 2, 3, 4), nrow=2)
> y
      [,1] [,2]
[1,]    1    3
[2,]    2    4
>
```


Creating Matrices (Cont.)

- When print out **y**, R shows its notation for rows and columns.

```
> y[, 2]  
[1] 3  4
```


Creating Matrices (Cont.)

- Another way to build **y** is to specify elements individually:

```
Console ~/ 
>
> y <- matrix(nrow=2, ncol=2)
> y[1,1] <- 1
> y[2,1] <- 2
> y[1,2] <- 3
> y[2,2] <- 4
> y
      [,1] [,2]
[1,]    1    3
[2,]    2    4
>
```



Creating Matrices (Cont.)

- Though internal storage of a matrix is in *column-major* order, you can set the **byrow** argument in **matrix()** to **true** to indicate that the data is coming in *row-major* order.

```
Console ~/   
> m <- matrix(c(1, 2, 3, 4, 5, 6), nrow=2, byrow=TRUE)  
> m  
      [,1] [,2] [,3]  
[1,]    1    2    3  
[2,]    4    5    6  
>
```

General Matrix Operations

- Performing Linear Algebra Operations on Matrices.

```
Console ~/ 
> # Mathematical matrix multiplication
> y
      [,1] [,2]
[1,]    1    3
[2,]    2    4
> y %*% y
      [,1] [,2]
[1,]    7   15
[2,]   10   22
>
> # Mathematical multiplication of matrix by scalar
> 3 * y
      [,1] [,2]
[1,]    3    9
[2,]    6   12
>
> # Mathematical matrix addition
> y + y
      [,1] [,2]
[1,]    2    6
[2,]    4    8
```

Matrix Indexing

Console ~/ ↩

```
> z <- matrix(c(1, 2, 3, 4, 1, 1, 0, 0, 1, 0, 1, 0), nrow=4)
```

```
> z
```

	[,1]	[,2]	[,3]
[1,]	1	1	1
[2,]	2	1	0
[3,]	3	0	1
[4,]	4	0	0

```
>
```

```
> z[, 2:3]
```

	[,1]	[,2]
[1,]	1	1
[2,]	1	0
[3,]	0	1
[4,]	0	0

```
>
```

Matrix Indexing (Cont.)

Console ~/ ↻

```
> y <- matrix(c(11, 12, 21, 22, 31, 32), nrow=3, byrow = T)
```

```
> y
```

```
      [,1] [,2]
[1,]   11   12
[2,]   21   22
[3,]   31   32
```

```
>
```

```
> y[2:3, ]
```

```
      [,1] [,2]
[1,]   21   22
[2,]   31   32
```

```
>
```

```
> y[2:3, 2]
```

```
[1] 22 32
```

```
>
```

Matrix Indexing (Cont.)

Console ~/ ↩

```
> y <- matrix(c(1, 2, 3, 4, 5, 6), nrow=3)
```

```
> y
```

	[,1]	[,2]
[1,]	1	4
[2,]	2	5
[3,]	3	6

```
>
```

```
> y[c(1, 3), ] <- matrix(c(1, 1, 8, 12), nrow=2)
```

```
> y
```

	[,1]	[,2]
[1,]	1	8
[2,]	2	5
[3,]	1	12

```
>
```


Matrix Indexing (Cont.)

Console ~/ ↻

```
> x <- matrix(nrow=3, ncol=3)
> y <- matrix(c(4, 5, 2, 3), nrow=2)
> y
      [,1] [,2]
[1,]    4    2
[2,]    5    3
>
> x[2:3, 2:3] <- y
> x
      [,1] [,2] [,3]
[1,]   NA   NA   NA
[2,]   NA    4    2
[3,]   NA    5    3
>
```

Matrix Indexing (Cont.)

- Negative subscripts, used with vectors to exclude certain elements, work the same way with matrices:

```
Console ~/   
> y <- matrix(c(1, 2, 3, 4, 5, 6), nrow=3)  
> y  
      [,1] [,2]  
[1,]    1    4  
[2,]    2    5  
[3,]    3    6  
>  
> y[-2, ]  
      [,1] [,2]  
[1,]    1    4  
[2,]    3    6  
>
```

Filtering on Matrices

- Filtering can be done with matrices, just as with vectors.
- You must be careful with the syntax, though.

```
Console C:/R Home/ ↗
> x <- matrix(c(1,2,3,2,3,4), nrow = 3)
> x
      [,1] [,2]
[1,]    1    2
[2,]    2    3
[3,]    3    4
>
> x[x[,2] >= 3, ]
      [,1] [,2]
[1,]    2    3
[2,]    3    4
```

```
Console C:/R Home/ ↗
>
> j <- x[,2] >= 3
> j
[1] FALSE  TRUE  TRUE
>
> x[j, ]
      [,1] [,2]
[1,]    2    3
[2,]    3    4
```


Filtering on Matrices (Cont.)

- Look!

```
Console C:/R Home/ ↗
> x <- matrix(c(1, 2, 3, 2, 3, 4), nrow=3)
> x
      [,1] [,2]
[1,]    1    2
[2,]    2    3
[3,]    3    4
>
> x[x[, 2] >= 3,]
      [,1] [,2]
[1,]    2    3
[2,]    3    4
```

- The object **x[,2]** is a vector.
- The operator **>=** compares two vectors.
- The number **3** was recycled to a vector of 3s.

Filtering on Matrices (Cont.)

- Look!

```
Console C:/R Home/ ↵  
> x  
      [,1] [,2]  
[1,]    1    2  
[2,]    2    3  
[3,]    3    4  
>  
> z <- c(5, 12, 13)  
> x[z %% 2 == 1, ]  
      [,1] [,2]  
[1,]    1    2  
[2,]    3    4
```

- The expression **`z %% 2 == 1`** tests each element of `z` for being an odd number, thus yielding (**TRUE, FALSE, TRUE**).
- As a result, we extracted the first and third rows of **`x`**.

Filtering on Matrices (Cont.)

- Look!

```
Console C:/R Home/ ↗
> m <- matrix(c(1,2,3,4,5,6), nrow=3)
> m
      [,1] [,2]
[1,]    1    4
[2,]    2    5
[3,]    3    6
>
> m[m[,1] > 1 & m[,2] > 5, ]
[1] 3 6
```

- First, the expression **`m[,1] > 1`** compares each element of the first column of m to 1 and returns (**`FALSE,TRUE,TRUE`**).
- The second expression, **`m[,2] > 5`**, similarly returns (**`FALSE,FALSE,TRUE`**).
- We then take the logical **AND** of (**`FALSE,TRUE,TRUE`**) and (**`FALSE,FALSE,TRUE`**), yielding (**`FALSE,FALSE,TRUE`**).

Filtering on Matrices (Cont.)

- Look!

Console C:/R Home/ ↗

```
> m <- matrix(c(5, 2, 9, -1, 10, 11), nrow=3)
```

```
> m
```

	[, 1]	[, 2]
[1,]	5	-1
[2,]	2	10
[3,]	9	11

```
>
```

```
> which(m > 2)
```

```
[1] 1 3 5 6
```

Applying Functions to Matrix Rows and Columns

- One of the most famous and most used features of R is the ***apply()** family of functions.
- That is, **apply()**, **tapply()**, and **lapply()**.
- Here, we'll look at **apply()**, which instructs R to call a user-specified function on each of the rows or each of the columns of a matrix.

Using the apply() Function

- Is the general form of apply for matrices:

`apply(m, dimcode, f, fargs)`

- The arguments:

- `m` is the matrix.
- `dimcode` is the dimension, equal to 1 if the function applies to rows or 2 for columns.
- `f` is the function to be applied.
- `fargs` is an optional set of arguments to be supplied to `f`.

Using the apply() Function (Cont.)

Console C:/R Home/ ↗

```
>  
> z <- matrix(c(1, 2, 3, 4, 5, 6), nrow=3)  
> z  
      [,1] [,2]  
[1,]    1    4  
[2,]    2    5  
[3,]    3    6  
>  
> apply(z, 2, mean)  
[1] 2 5
```

Using the apply() Function (Cont.)

• Look!

```
Console C:/R Home/ ↗
> z
      [,1] [,2]
[1,]    1    4
[2,]    2    5
[3,]    3    6
>
> f <- function(x) x/c(2,8)
> y <- apply(z, 1, f)
> y
      [,1] [,2] [,3]
[1,]  0.5 1.000 1.50
[2,]  0.5 0.625 0.75
```

- `f()` function divides a two-element vector by the vector `(2, 8)`.
- Recycling would be used if `x` had a length longer than `2`.
- The call to `apply()` asks R to call `f()` on each of the rows of `z`.
- The first such row is `(1, 4)`, so in the call to `f()`, the actual argument corresponding to the formal argument `x` is `(1, 4)`.
- Thus, R computes the value of `(1, 4) / (2, 8)`, which in R's element-wise vector arithmetic is `(0.5, 0.5)`.

Using the apply() Function (Cont.)

- If the function to be applied returns a vector of **k** components, then the result of **apply()** will have **k** rows.
- Can use the matrix transpose function **t()** to change it if necessary.

```
Console C:/R Home/ ↗
> z
      [,1] [,2]
[1,]    1    4
[2,]    2    5
[3,]    3    6
>
> f
function(x) x/c(2,8)
<bytecode: 0x000000000037ae178>
>
> t(apply(z, 1, f))
      [,1] [,2]
[1,]  0.5 0.500
[2,]  1.0 0.625
[3,]  1.5 0.750
```

Using the apply() Function (Cont.)

Console C:/R Home/ ↗

```
> copymaj <- function(rw, d){  
+   maj <- sum(rw[1:d]) / d  
+   return(if(maj > 0.5) 1 else 0)  
+ }  
>  
> x <- matrix(c(1,1,1,0,0,1,0,1,1,1,0,1,1,1,1,1,0,0,1,0), nrow=4)  
> x  
      [,1] [,2] [,3] [,4] [,5]  
[1,]    1    0    1    1    0  
[2,]    1    1    1    1    0  
[3,]    1    0    0    1    1  
[4,]    0    1    1    1    0  
>  
> apply(x, 1, copymaj, 3)  
[1] 1 1 0 1  
>  
> apply(x, 1, copymaj, 2)  
[1] 0 1 0 0
```

Adding and Deleting Matrix Rows and Columns

- Technically, matrices are of fixed length and dimensions, so we cannot add or delete rows or columns.
- However, matrices can be reassigned.
- Can achieve the same effect as if we had directly done additions or deletions.

Changing the Size of a Matrix

• Look!

```
Console C:/R Home/ ↗
> x <- c(12, 5, 13, 16, 8)
> x
[1] 12  5 13 16  8
>
> x <- c(x, 20)    # Append 20
> x
[1] 12  5 13 16  8 20
>
> x <- c(x[1:3], 20, x[4:6])  # Insert 20
> x
[1] 12  5 13 20 16  8 20
>
> x <- x[-2:-4]    # Delete elements 2 through 4
> x
[1] 12 16  8 20
```

- **x** is originally of length **5**, which we extend to **6** via concatenation and then reassignment.
- We didn't literally change the length of **x** but instead created a new vector from **x** and then assigned **x** to that new vector.


Changing the Size of a Matrix (Cont.)

- The **`rbind()`** (*row bind*) and **`cbind()`** (*column bind*) functions let you add rows or columns to a matrix.

```
Console C:/R Home/ ↗
> one <- c(1, 1, 1, 1)
> one
[1] 1 1 1 1
>
> z <- matrix(c(1,2,3,4,1,1,0,0,1,0,1,0), nrow=4)
> z
      [,1] [,2] [,3]
[1,]    1    1    1
[2,]    2    1    0
[3,]    3    0    1
[4,]    4    0    0
>
> cbind(one, z)
      one
[1,]    1 1 1 1
[2,]    1 2 1 0
[3,]    1 3 0 1
[4,]    1 4 0 0
```

Changing the Size of a Matrix (Cont.)

- **cbind()** creates a new matrix by combining a column of 1s with the columns of **z**.

```
Console C:/R Home/   
> one  
[1] 1 1 1 1  
>  
> z  
      [,1] [,2] [,3]  
[1,]    1    1    1  
[2,]    2    1    0  
[3,]    3    0    1  
[4,]    4    0    0  
>  
> cbind(1, z)  
      [,1] [,2] [,3] [,4]  
[1,]    1    1    1    1  
[2,]    1    2    1    0  
[3,]    1    3    0    1  
[4,]    1    4    0    0
```

Changing the Size of a Matrix (Cont.)

- Can also use the **rbind()** and **cbind()** functions as a quick way to create small matrices.

```
Console C:/R Home/ ↵  
> q <- cbind(c(1,2), c(3, 4))  
> q  
      [,1] [,2]  
[1, ]    1    3  
[2, ]    2    4
```

Changing the Size of a Matrix (Cont.)

- Can delete rows or columns by reassignment.

```
Console C:/R Home/ ↗  
> m <- matrix(1:6, nrow=3)  
> m  
      [,1] [,2]  
[1,]    1    4  
[2,]    2    5  
[3,]    3    6  
>  
> m <- m[c(1,3), ]  
> m  
      [,1] [,2]  
[1,]    1    4  
[2,]    3    6
```


More on the Vector/Matrix Distinction

```
Console C:/R Home/ ↗
> z <- matrix(1:8, nrow=4)
> z
      [,1] [,2]
[1,]    1    5
[2,]    2    6
[3,]    3    7
[4,]    4    8
>
> length(z)
[1] 8
>
> class(z)
[1] "matrix"
>
> attributes(z)
$dim
[1] 4 2
```

- As **z** is still a vector, we can query its **length**.
- But as a matrix, **z** is a bit more than a vector.
- Most of R consists of **S3** classes.
- S3 components are denoted by dollar signs (**\$**).
- The matrix class has one attribute, named **dim**, which is a vector containing the numbers of rows and columns in the matrix.

More on the Vector/Matrix Distinction (Cont.)

- You can also obtain dim via the `dim()` function.

```
> dim(z)
[1] 4 2
```

- The numbers of rows and columns are obtainable individually via the `nrow()` and `ncol()` functions.

```
> nrow(z)
[1] 4
>
> ncol(z)
[1] 2
```

Avoiding Unintended Dimension Reduction

- In the world of statistics, dimension reduction is a good thing, with many statistical procedures aimed to do it well.
- If we are working with, say, 10 variables and can reduce that number to 3 that still capture the essence of our data, we're happy.
- However, in R, something else might merit the name *dimension reduction* that we may sometimes wish to avoid.

Avoiding Unintended Dimension Reduction

Console C:/R Home/ ↻

```
>
> z <- matrix(1:8, nrow=4)
> z
      [,1] [,2]
[1,]    1    5
[2,]    2    6
[3,]    3    7
[4,]    4    8
>
> r <- z[2,]
> r
[1] 2 6
```

- **r**, it's a vector format, not a matrix format.
- In other words, **r** is a vector of length **2**, rather than a **1-by-2** matrix.

Avoiding Unintended Dimension Reduction (Cont.)

```
> z
      [,1] [,2]
[1,]    1    5
[2,]    2    6
[3,]    3    7
[4,]    4    8
>
> attributes(z)
$dim
[1] 4 2

> r
[1] 2 6
>
> attributes(r)
NULL
>
> str(z)
int [1:4, 1:2] 1 2 3 4 5 6 7 8
>
> str(r)
int [1:2] 2 6
```

- R informs us that **z** has row and column numbers, while **r** does not.
- Similarly, **str()** tells us that **z** has indices ranging in **1:4** and **1:2**, for rows and columns, while **r**'s indices simply range in **1:2**.
- No doubt about it—**r** is a vector, not a matrix.


Avoiding Unintended Dimension Reduction (Cont.)

```
Console C:/R Home/ ↗
> z
      [,1] [,2]
[1,]    1    5
[2,]    2    6
[3,]    3    7
[4,]    4    8
>
> r <- z[2,, drop=FALSE]
> r
      [,1] [,2]
[1,]    2    6
>
> dim(r)
[1] 1 2
```

- Fortunately, R has a way to suppress this dimension reduction.
- The **drop** argument.
- Here's an example, using the matrix **z**.
- Now **r** is a **1-by-2** matrix, not a two-element vector.

Avoiding Unintended Dimension Reduction (Cont.)

- `[` is actually a function, just as is the case for operators like `+`.

```
Console C:/R Home/   
>  
> z  
      [,1] [,2]  
[1, ]    1    5  
[2, ]    2    6  
[3, ]    3    7  
[4, ]    4    8  
>  
> z[3,2]  
[1] 7  
>  
> "[(z, 3, 2)  
[1] 7  
>
```

Avoiding Unintended Dimension Reduction (Cont.)

Console C:/R Home/ ↗

```
> u <- c(1,2,3)
>
> u
[1] 1 2 3
>
> v <- as.matrix(u)
>
> attributes(u)
NULL
>
> attributes(v)
$dim
[1] 3 1
```

- If you have a vector that you wish to be treated as a matrix, you can use the **as.matrix()** function.

Naming Matrix Rows and Columns

```
Console C:/R Home/ ↗
> z <- matrix(1:4, nrow=2)
> z
      [,1] [,2]
[1,]    1    3
[2,]    2    4
>
> colnames(z)
NULL
>
> colnames(z) <- c("a", "b")
>
> z
      a b
[1,] 1 3
[2,] 2 4
>
> colnames(z)
[1] "a" "b"
>
> z[, "a"]
[1] 1 2
```

- The natural way to refer to rows and columns in a matrix is via the row and column numbers.
- However, you can also give names to these entities.
- These names can then be used to reference specific columns.
- The function **rownames()** works similarly.

Higher-Dimensional Arrays

- The matrix is then a two-dimensional data structure.
- But suppose also have data taken at different times, one data point per person per variable per time.
- Time then becomes the third dimension, in addition to rows and columns.
- In R, such data sets are called *arrays*.

Higher-Dimensional Arrays (Cont.)

- Here's the data for the first test.

```
Console C:/R Home/ ↗  
> firsttest <- matrix(c(46, 21, 50, 30, 25, 50), nrow=3)  
> firsttest  
      [,1] [,2]  
[1,]   46  30  
[2,]   21  25  
[3,]   50  50  
>
```

- Student 1 had scores of 46 and 30 on the first test, student 2 scored 21 and 25, and so on.

Higher-Dimensional Arrays (Cont.)

- Here's the data for the second test.

```
> secondtest <- matrix(c(46, 41, 50, 43, 35, 50), nrow=3)
> secondtest
      [,1] [,2]
[1,]   46   43
[2,]   41   35
[3,]   50   50
>
```

Higher-Dimensional Arrays (Cont.)

- Now let's put both tests into one data structure, which we'll name **tests**.
- We'll arrange it to have two *layers*—one layer per test—with three rows and two columns within each layer.
- We'll store **firsttest** in the first layer and **secondtest** in the second.
- In layer 1, there will be three rows for the three students' scores on the first test, with two columns per row for the two portions of a test.
- We use R's array function to create the data structure:

```
> tests <- array(data=c(firsttest,secondtest),dim=c(3,2,2))
```

Higher-Dimensional Arrays (Cont.)

- In the argument `dim=c(3,2,2)`, we are specifying two layers (this is the second 2), each consisting of three rows and two columns.

```
Console C:/R Home/ ↗
> firsttest
      [,1] [,2]
[1,]   46  30
[2,]   21  25
[3,]   50  50
>
> secondtest
      [,1] [,2]
[1,]   46  43
[2,]   41  35
[3,]   50  50
>
> tests <- array(data=c(firsttest,secondtest),dim=c(3,2,2))
>
> attributes(tests)
$dim
[1] 3 2 2
```

Higher-Dimensional Arrays (Cont.)

Console C:/R Home/ ↗

```
>
> tests
, , 1

      [,1] [,2]
[1,]   46   30
[2,]   21   25
[3,]   50   50

, , 2

      [,1] [,2]
[1,]   46   43
[2,]   41   35
[3,]   50   50

>
> tests[3,2,1]
[1] 50
```

- Each element of **tests** now has three subscripts, rather than two as in the matrix case.
- The first subscript corresponds to the first element in the **\$dim** vector, the second subscript corresponds to the second element in the vector, and so on.