

Data Frames

Bok, Jong Soon
javaexpert@nate.com
www.javaexpert.info

Data Frames

- On an intuitive level, a data frame is like a matrix, with a two-dimensional rows-and-columns structure.
- However, it differs from a matrix in that each column may have a different mode.
- For instance, one column may consist of numbers, and another column might have character strings.
- In this sense, just as lists are the heterogeneous analogs of vectors in one dimension, data frames are the heterogeneous analogs of matrices for two-dimensional data.

Creating Data Frames

- Let's begin.

Console C:/R Home/ ↗

```
> kids <- c("Jack", "Jill")
> ages <- c(12, 10)
> d <- data.frame(kids, ages, stringsAsFactors = FALSE)
> d
```

	kids	ages
1	Jack	12
2	Jill	10

Creating Data Frames (Cont.)

- The first two arguments in the call to `data.frame()` are clear.
- Produce a data frame from our two vectors: kids and ages.
- If the named argument `stringsAsFactors` is not specified, then by default, `stringsAsFactors` will be `TRUE`.
- This means that if we create a data frame from a character vector, kids \rightarrow R will convert that vector to a factor.
- Because our work with character data will typically be with vectors rather than factors, we'll set `stringsAsFactors` to `FALSE`.

Accessing Data Frames

```
Console C:/R Home/ ↵  
> d  
  kids ages  
1 Jack   12  
2 Jill   10  
>  
> d[[1]]  
[1] "Jack" "Jill"  
>  
> d$kids  
[1] "Jack" "Jill"
```

- Since **d** is a list, can access it as such via component index values or component names.

Accessing Data Frames (Cont.)

- But Can treat it in a matrix-like fashion as well.
- For example, can view column 1.

```
Console C:/R Home/ ↵  
> d  
  kids ages  
1 Jack   12  
2 Jill   10  
>  
> d[,1]  
[1] "Jack" "Jill"
```

Accessing Data Frames (Cont.)

- This matrix-like quality is also seen when take **d** apart using **str()**.

```
Console C:/R Home/ ↗
> d
  kids ages
1 Jack   12
2 Jill   10
>
> str(d)
'data.frame':   2 obs. of  2 variables:
 $ kids: chr  "Jack" "Jill"
 $ ages: num  12 10
```

Accessing Data Frames (Cont.)

- Consider three ways to access the first column of our data frame above
 - `d[[1]]`
 - `d[,1]`
 - `d$kids`
- Of these, the third would generally be considered to be clearer and, more importantly, safer than the first two.
- This better identifies the column and makes it less likely that you will reference the wrong column.

Other Matrix-Like Operations

- Various matrix operations also apply to data frames.
- Most notably and usefully, can do filtering to extract various subdata frames of interest.

Extracting Subdata Frames

- A data frame can be viewed in row-and-column terms.
- In particular, we can extract subdata frames by rows or columns.

Extracting Subdata Frames (Cont.)

ExamsQuiz.txt - Notepad

File Edit Format View Help

"Exam 1"	"Exam 2"	Quiz
2.0	3.3	4.0
3.3	2.0	3.7
4.0	4.0	4.0
2.3	0.0	3.3
2.3	1.0	3.3
3.3	3.7	4.0

Extracting Subdata Frames (Cont.)

Console C:/R Home/ ↗

```
> examsquiz <- read.table(file = "ExamsQuiz.txt", header = TRUE)  
>
```

```
> head(examsquiz)
```

	Exam.1	Exam.2	Quiz
1	2.0	3.3	4.0
2	3.3	2.0	3.7
3	4.0	4.0	4.0
4	2.3	0.0	3.3
5	2.3	1.0	3.3
6	3.3	3.7	4.0

Extracting Subdata Frames (Cont.)

```
> examsquiz[2:5, ]
  Exam.1 Exam.2 Quiz
2    3.3     2  3.7
3    4.0     4  4.0
4    2.3     0  3.3
5    2.3     1  3.3
>
> examsquiz[2:5, 2]
[1] 2 4 0 1
>
> class(examsquiz[2:5, 2])
[1] "numeric"
>
> examsquiz[2:5, 2, drop=FALSE]
  Exam.2
2      2
3      4
4      0
5      1
>
> class(examsquiz[2:5, 2, drop=FALSE])
[1] "data.frame"
```

- Since `examsquiz[2:5, 2]` is a vector, R created a vector instead of another data frame.
- By specifying `drop=FALSE`, can keep it as a (one-column) data frame.

Extracting Subdata Frames (Cont.)

- We can also do filtering. Here's how to extract the subframe of all students whose first exam score was at least 3.8.

Console C:/R Home/ ↗

```
> examsquiz[examsquiz$Exam.1 >= 3.8, ]
```

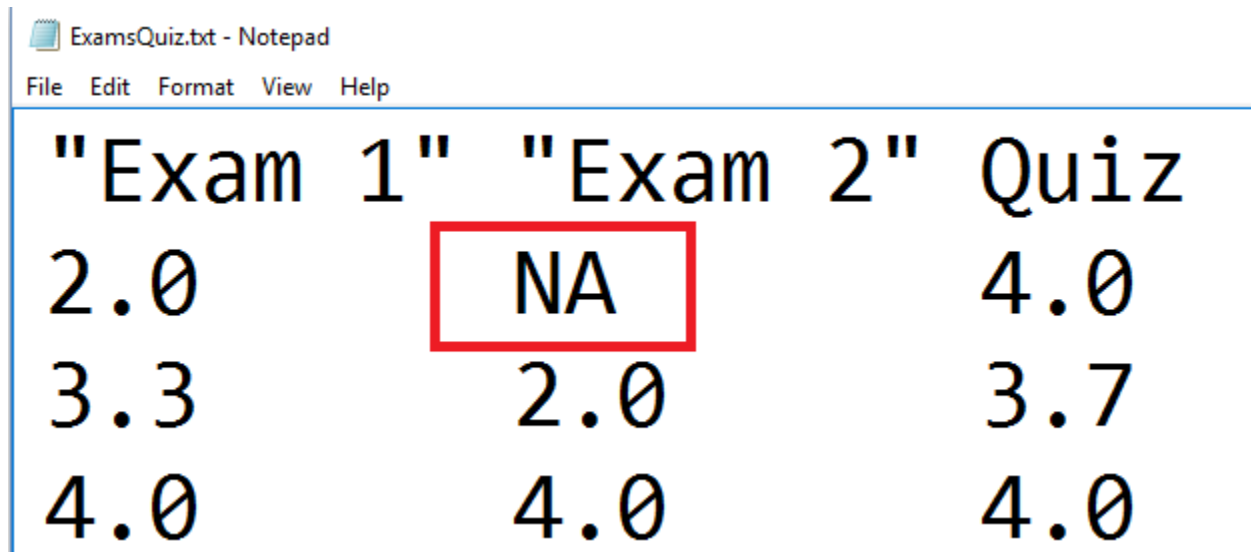
```
Exam.1 Exam.2 Quiz
```

```
3      4      4      4
```

```
>
```

More on Treatment of NA Values

- Suppose the second exam score for the first student had been missing.
- Then we would have typed the following into that line when we were preparing the data file.



A screenshot of a Notepad window titled "ExamsQuiz.txt - Notepad". The window displays a table with three columns: "Exam 1", "Exam 2", and "Quiz". The first row shows scores of 2.0, NA, and 4.0. The second row shows 3.3, 2.0, and 3.7. The third row shows 4.0, 4.0, and 4.0. The "NA" value in the first row is highlighted with a red rectangular box.

"Exam 1"	"Exam 2"	Quiz
2.0	NA	4.0
3.3	2.0	3.7
4.0	4.0	4.0

More on Treatment of NA Values (Cont.)

- In any subsequent statistical analyses, R would do its best to cope with the missing data.
- However, in some situations, we need to set the option **na.rm=TRUE**, explicitly telling R to ignore **NA** values.
- For instance, with the missing exam score, calculating the mean score on exam 2 by calling R's **mean()** function would skip that first student in finding the mean.
- Otherwise, R would just report **NA** for the mean.

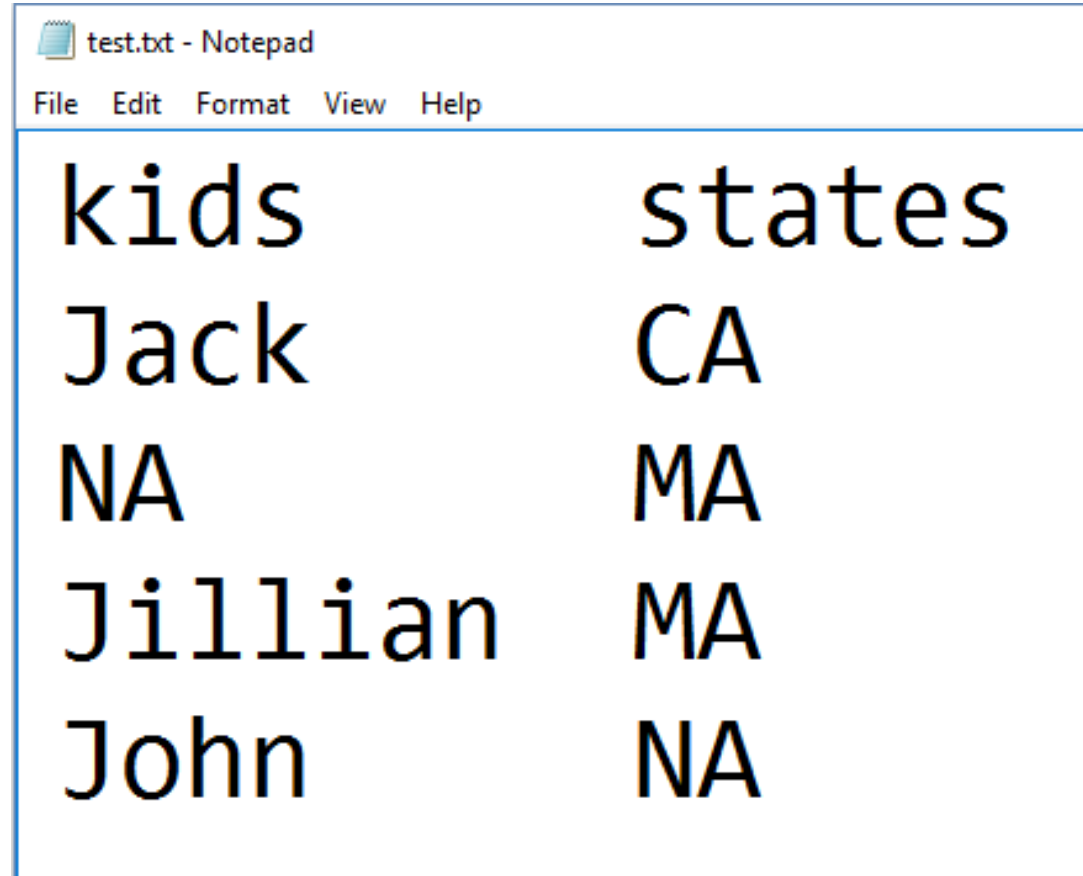
```
Console C:/R Home/ ↗  
> x <- c(2, NA, 4)  
> mean(x)  
[1] NA  
>  
> mean(x, na.rm = TRUE)  
[1] 3
```


More on Treatment of NA Values (Cont.)

- **subset()** function, which saves you the trouble of specifying **na.rm=TRUE**.
- Can apply it in data frames for row selection.
- The column names are taken in the context of the given data frame.

```
Console C:/R Home/ ↵
> examsquiz <- read.table(file = "ExamsQuiz.txt", header = TRUE)
>
> examsquiz[examsquiz$Exam.1 >= 3.8, ]
  Exam.1 Exam.2 Quiz
3      4      4    4
>
> subset(examsquiz, Exam.1 >= 3.8)
  Exam.1 Exam.2 Quiz
3      4      4    4
```

More on Treatment of NA Values (Cont.)



A screenshot of a Notepad window titled "test.txt - Notepad". The window contains a table with two columns: "kids" and "states". The table has five rows of data. The first row shows "kids" and "states". The second row shows "Jack" and "CA". The third row shows "NA" and "MA". The fourth row shows "Jillian" and "MA". The fifth row shows "John" and "NA".

kids	states
Jack	CA
NA	MA
Jillian	MA
John	NA

More on Treatment of NA Values (Cont.)

```
Console C:/R Home/ ↗
> d4 <- read.table(file = "test.txt", header = TRUE)
Warning message:
In read.table(file = "test.txt", header = TRUE) :
  incomplete final line found by readTableHeader on 'test.txt'
>
> d4
      kids states
1    Jack    CA
2   <NA>    MA
3 Jillian    MA
4    John  <NA>
>
> complete.cases(d4)
[1]  TRUE FALSE  TRUE FALSE
>
> d5 <- d4[complete.cases(d4), ]
> d5
      kids states
1    Jack    CA
3 Jillian    MA
```

- In some cases, we may wish to rid our data frame of any observation that has at least one **NA** value.
- A handy function for this purpose is **complete.cases()**.

Using the rbind() and cbind() Functions and Alternatives

- Can use **cbind()** to add a new column that has the same length as the existing columns.
- In using **rbind()** to add a row, the added row is typically in the form of another data frame or list.

```
Console C:/R Home/ ↗
> kids <- c("Jack", "Jill")
> ages <- c(12, 10)
> d <- data.frame(kids, ages, stringsAsFactors = FALSE)
> d
  kids ages
1 Jack  12
2 Jill  10
>
> rbind(d, list("Laura", 19))
  kids ages
1 Jack  12
2 Jill  10
3 Laura 19
```

Using the rbind() and cbind() Functions and Alternatives (Cont.)

- Can also create new columns from old ones.
- For instance, can add a variable that is the difference between exams 1 and 2.

Using the rbind() and cbind() Functions and Alternatives (Cont.)

```
Console C:/R Home/ ↗  
> examsquiz  
  Exam.1 Exam.2 Quiz  
1    2.0    3.3  4.0  
2    3.3    2.0  3.7  
3    4.0    4.0  4.0  
4    2.3    0.0  3.3  
5    2.3    1.0  3.3  
6    3.3    3.7  4.0  
>  
> eq <- cbind(examsquiz, examsquiz$Exam.2 - examsquiz$Exam.1)  
> class(eq)  
[1] "data.frame"  
>  
> head(eq)  
  Exam.1 Exam.2 Quiz examsquiz$Exam.2 - examsquiz$Exam.1  
1    2.0    3.3  4.0                                1.3  
2    3.3    2.0  3.7                               -1.3  
3    4.0    4.0  4.0                                 0.0  
4    2.3    0.0  3.3                               -2.3  
5    2.3    1.0  3.3                               -1.3  
6    3.3    3.7  4.0                                 0.4
```

Using the rbind() and cbind() Functions and Alternatives (Cont.)

- The new name is rather unwieldy.
- It's long, and it has embedded blanks.

Console C:/R Home/ ↗

```
> examsquiz$ExamDiff <- examsquiz$Exam.2 - examsquiz$Exam.1
```

```
> head(examsquiz)
```

	Exam.1	Exam.2	Quiz	ExamDiff
1	2.0	3.3	4.0	1.3
2	3.3	2.0	3.7	-1.3
3	4.0	4.0	4.0	0.0
4	2.3	0.0	3.3	-2.3
5	2.3	1.0	3.3	-1.3
6	3.3	3.7	4.0	0.4

Using the rbind() and cbind() Functions and Alternatives (Cont.)

Console C:/R Home/ ↗

```
> d
  kids ages
1 Jack   12
2 Jill   10
>
> d$one <- 1
>
> d
  kids ages one
1 Jack   12   1
2 Jill   10   1
```

- Can even exploit recycling to add a column that is of a different length than those in the data frame.

Applying apply()

- You can use `apply()` on data frames, if the columns are all of the same type.
- For instance, we can find the maximum grade for each student.

```
Console C:/R Home/ ↗
> examsquiz
  Exam.1 Exam.2 Quiz
1    2.0    3.3  4.0
2    3.3    2.0  3.7
3    4.0    4.0  4.0
4    2.3    0.0  3.3
5    2.3    1.0  3.3
6    3.3    3.7  4.0
>
> apply(examsquiz, 1, max)
[1] 4.0 3.7 4.0 3.3 3.3 4.0
```

Merging Data Frames

- In the relational database world, one of the most important operations is that of a *join*.
- Join two tables can be combined according to the values of a common variable.
- In R, two data frames can be similarly combined using the `merge()` function.

Merging Data Frames (Cont.)

- The simplest form is as follows:

`merge(x, y)`

- This merges data frames x and y.
- It assumes that the two data frames have one or more columns with names in common.

Merging Data Frames (Cont.)

"kids"	"states"
Jack	CA
Jill	MA
Jillian	MA
John	HI

```
Console C:/R Home/ ↗
> d1 <- read.table(file = "d1.txt", header = TRUE)
Warning message:
In read.table(file = "d1.txt", header = TRUE) :
  incomplete final line found by readTableHeader on 'd1.txt'
>
> ages <- c(10, 7, 12)
> kids <- c("Jill", "Lillian", "Jack")
> d2 <- data.frame(ages, kids, stringsAsFactors = FALSE)
>
> d1
      kids states
1   Jack     CA
2   Jill     MA
3 Jillian     MA
4   John     HI
>
> d2
  ages   kids
1   10   Jill
2    7 Lillian
3   12   Jack
```

Merging Data Frames (Cont.)

- The two data frames have the variable kids in common.
- R found the rows in which this variable had the same value of kids in both data frames(the ones for Jack and Jill).
- It then created a data frame with corresponding rows and with columns taken from data frames (kids, states, and ages).

```
> d <- merge(d1, d2)
```

```
> d
```

	kids	states	ages
1	Jack	CA	12
2	Jill	MA	10

Merging Data Frames (Cont.)

Console C:/R Home/ ↗

```
> d3 <- read.table(file = "d3.txt", header = TRUE)
```

Warning message:

```
In read.table(file = "d3.txt", header = TRUE) :  
  incomplete final line found by readTableHeader on 'd3.txt'
```

```
>
```

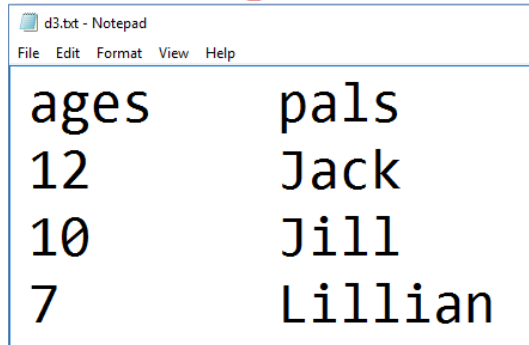
```
> d3
```

	ages	pals
1	12	Jack
2	10	Jill
3	7	Lillian

```
>
```

```
> merge(d1, d3, by.x = "kids", by.y = "pals")
```

	kids	states	ages
1	Jack	CA	12
2	Jill	MA	10



ages	pals
12	Jack
10	Jill
7	Lillian

- The **merge()** function has named arguments **by.x** and **by.y**, which handle cases in which variables have similar information but different names in the two data frames

Merging Data Frames (Cont.)

Console C:/R Home/ ↗

```
> d1
  kids states
1  Jack   CA
2  Jill   MA
3 Jillian MA
4  John   HI
>
> d2a <- rbind(d2, list(15, "Jill"))
> d2a
  ages  kids
1   10  Jill
2    7 Lillian
3   12  Jack
4   15  Jill
>
> merge(d1, d2a)
  kids states ages
1 Jack    CA   12
2 Jill    MA   10
3 Jill    MA   15
```

- Duplicate matches will appear in full in the result, possibly in undesirable ways.
- There are *two* Jills in d2a.
- There is a Jill in d1 who lives in MA and another Jill with unknown residence.
- In previous example, `merge(d1, d2)`, there was only one Jill, who was presumed to be the same person in both data frames.
- But here, in the call `merge(d1, d2a)`, it may have been the case that only one of the **Jills** was a MS resident.

Applying Functions to Data Frames

- As with lists, you can use the `lapply` and `sapply` functions with data frames.

Using lapply() and sapply() on Data Frames

```
Console C:/R Home/ ↗
> kids <- c("Jack", "Jill")
> ages <- c(12, 10)
> d <- data.frame(kids, ages, stringsAsFactors = FALSE)
>
> d
  kids ages
1 Jack   12
2 Jill   10
>
> d1 <- lapply(d, sort)
> d1
$kids
[1] "Jack" "Jill"

$ages
[1] 10 12
```

- Data frames are special cases of lists, with the list components consisting of the data frame's columns.
- Thus, if you call **lapply()** on a data frame with a specified function **f()**, then **f()** will be called on each of the frame's columns, with the return values placed in a list.

Using lapply() and sapply() on Data Frames (Cont.)

- Note that dl is just a list, not a data frame. We could coerce it to a data frame, like this.

```
Console C:/R Home/ ↗  
> as.data.frame(dl)  
  kids ages  
1 Jack   10  
2 Jill   12
```

- But this would make no sense, as the correspondence between names and ages has been lost.
- **Jack**, for instance, is now listed as **10** years old instead of **12**.